

CSCI-GA.2565 — Final Project

your NetID here

Version 1.0

Instructions.

- **Deliverables.**

1. Grading is based on a **poster presentation**
2. By **noon EDT on Wednesday, April 30**, you must complete a perfunctory handin on gradescope which lists your group members and confirms that you can attend the final poster session.

- **Group sizes.** You may work in groups up to 5.

- **Consulting LLMs and friends.** As usual, you must list all sources, which means answering questions about them in the oral presentation.

- **Evaluation.** Show up to the poster session, answer questions about your poster, visit and discuss the posters made by other groups.

Version history.

- 1.0. Initial version.

Suggestions.

1. **Extend the pico-llm project.** For instance:
 - Complete the LLM training pipeline: e.g., fine-tune on question-answer pairs, and/or perform RLHF/DPO.
 - Implement the DeepSeek R1-0 recipe and obtain a reasoning model.
 - Expand or initiate the interpretability analysis suggested in the pico-llm writeup.
 - Anything else that seems reasonable and fun to you.
2. Perform a literature survey. (*This writeup will eventually be expanded with some suggestions.*)
3. Any other ML project which is of clear relevance and overlap with course topics.